


N O R T H L I G H T



Northlight Weather API

Contents

Northlight Weather API	3
Getting started	7
Endpoint reference	12
Rate limits	19
Error codes	24
Release notes	27

Northlight Weather API v1

Northlight serves current conditions, short-range forecasts and observation history for a network of weather stations along the Nordic coast. This handbook is one set of Markdown documents that cudoc publishes three ways: as the site you may be reading, as a bound PDF with a cover and a contents page, and as a Word file with real styles. Every figure below is drawn from the reference pages, so a correction there reaches this page too.

How this handbook is built

The site, the PDF and the Word file come from one `cudoc-export build` run and one set of design tokens. The configuration is a single file, `showcase.config.mjs`, beside the documents, and the summary table further down is assembled from the endpoint reference at build time.

What the API provides

Every station reports temperature, wind, humidity, pressure and precipitation on a fixed schedule, and the forecast model runs once an hour for every station. The API exposes that data as four data sets, each behind one or two endpoints.

Data set	Cadence	Retention	Endpoints
Current conditions	Every 10 minutes	Latest reading only	Current conditions
Forecasts	Hourly model run	48 hours and 7 days ahead	Hourly forecast , Daily forecast
Observation history	Every 10 minutes	Depends on the plan	Observation history
Station metadata	On change	Indefinite	Station directory , Station detail

Forecast horizons and history retention depend on the plan; see [Plans compared](#).

Endpoints at a glance

Each row is one section of the endpoint reference. The method and the path are read from the table at the top of that section, and the description is its first paragraph, so this table cannot drift from the reference.

Endpoint	Method	Path	What it returns
Current conditions	GET	<code>/v1/weather/current</code>	Returns the latest observation from one station.

Endpoint	Method	Path	What it returns
Hourly forecast	GET	/v1/forecast/hourly	Returns up to 48 hourly forecast points for one station.
Daily forecast	GET	/v1/forecast/daily	Returns one forecast summary per calendar day, up to seven days ahead.
Observation history	GET	/v1/weather/history	Returns past observations from one station in a time range, oldest first.
Station directory	GET	/v1/stations	Lists the stations a key may read, with their location and status.
Station detail	GET	/v1/stations/{code}	Returns one station's metadata, including its sensors and time zone.
Weather alerts	GET	/v1/alerts	Lists active severe-weather alerts for a region, relayed from the national service.

Concepts

Stations and station codes

A station is one physical site with a fixed set of sensors. Every station has a code of the form **SITE-NN** : a three-letter site code and a two-digit sequence number, so a site with several masts has **BG0-01** , **BG0-02** and so on. Codes never change and are never reused; a station that stops reporting for good is marked **retired** and keeps its code.

Code	Site	Elevation	Status
OSL-01	Oslo	94 m	active
BG0-01	Bergen	12 m	active
BG0-02	Bergen	320 m	maintenance
TRD-01	Trondheim	8 m	active
TOS-01	Tromsø	100 m	active
SVG-01	Stavanger	4 m	retired

The list above is illustrative; the [station directory](#) is the source of truth for the stations a key may read.

Observations and forecasts

An observation is a reading a station actually made, stamped with the instant it was made. A forecast point is the model's expectation for a future instant, stamped with that instant and with the model run that produced it. The two are never mixed in one response: the current-conditions and history endpoints return observations, and the forecast endpoints return forecast points.

Units

Every numeric field is returned in metric units unless the request asks for `units=imperial`. Field names do not change with the unit system; only the values and the `units` object in the response do, so a client can always read the unit next to the value instead of assuming it.

Field	Metric	Imperial
temperature	°C	°F
windSpeed	m/s	mph
pressure	hPa	inHg
precipitation	mm	in
visibility	km	mi

Timestamps

Every instant is ISO 8601 in UTC with a `Z` suffix: `observedAt` on an observation, `validAt` and `issuedAt` on a forecast, `lastObservedAt` on a station. Converting to local time is the client's job; a station's IANA time zone is in its [detail record](#) so the conversion can be exact.

Versioning

The path prefix `/v1` names the version. Additive changes, such as a new field, a new optional parameter or a new endpoint, ship within `/v1` without notice, so a client must ignore fields it does not know. Any change to an existing shape is announced in the [release notes](#) at least ninety days before it takes effect, and the previous behaviour stays available for that period.

Conventions

The handbook uses a small set of marks, and each means the same thing in the site, the PDF and the Word file.

Mark	Meaning
A New , Beta or Latest badge after a heading	The section was added in the latest release, is in beta and may change on thirty days' notice, or describes the newest release.
A Note or Tip box	Background or a shortcut; skipping it costs nothing.
An Important , Warning or Caution box	Something that costs a key, a quota or data if ignored.
A Ready box	A checkpoint listing what a client needs before it moves on.
<code>code</code> in running text	A literal parameter, field, header or value.

Request and response samples are given as `curl` , Node.js and Python, and every JSON sample is a real response for the station `OSL-01` .

Where to start


1. Read [Getting started](#) for a key, a first request and a polling schedule.
2. Keep [Rate limits](#) open while you design that schedule; the request costs table tells you what each call weighs.
3. Handle every code in [Error codes](#) before you switch to a live key.
4. Check the [release notes](#) before upgrading a client.

Reading this on paper?

External links print their address after the text, and links between documents jump inside the bound file. The contents page at the front lists the page each document starts on.

Getting started

Four steps take you from nothing to a first forecast in production: a key, a request, a look at the response, and a polling schedule that respects the cache. The whole flow is plain HTTPS with a key in a header; there is no SDK to install, and any HTTP client will do.



GET /v1/weather/current?station=... → cached at the edge for 60 seconds → read from the store on a miss

Get a key

Keys are issued per project and scoped to the stations a project may read. A project is created in the Northlight console, and every key it issues carries the project's plan and its station list. Two kinds of key exist:

Kind	Prefix	Purpose
Live	nlk_live_	Reads real stations and counts against the project's plan.
Test	nlk_test_	Reads the synthetic station TEST-01 only, returns fixed data and is never rate limited.

Start with a test key. TEST-01 reports the same reading every ten minutes and a forecast that is always 12 °C and calm, so your parser can be checked against known values before a live key is involved.

Requirements

Step	Prerequisites
Register	- A project name - A billing contact - At least one station - Accepted terms - A callback URL

Step	Prerequisites
Rotate	- The current key - Owner role
Revoke	Owner role

Store the key in your secret manager and expose it to your service as an environment variable; the examples below read `NORTHLIGHT_KEY`.

Keep the key out of the browser

A key grants every scope its project holds. Call Northlight from your own backend and pass results on; never ship the key in client-side code, a mobile app or a public repository. If a key leaks, revoke it in the console and issue a new one; requests with a revoked key fail with `key_revoked`.

Make a request New

Send the key in the `Authorization` header and ask for a station by its code:

```
curl -H "Authorization: Bearer $NORTHLIGHT_KEY" \  
  "https://api.northlight.example/v1/weather/current?station=OSL-01"
```

```
{  
  "station": "OSL-01",  
  "observedAt": "2026-09-22T06:00:00Z",  
  "temperature": 11.4,  
  "windSpeed": 3.2,  
  "windDirection": 240,  
  "humidity": 87,  
  "pressure": 1009.6,  
  "precipitation": 0.0,  
  "conditions": "overcast",  
  "stale": false,  
  "units": {  
    "temperature": "C",  
    "windSpeed": "m/s",  
    "pressure": "hPa",  
    "precipitation": "mm"  
  }  
}
```

```
}  
}
```

The same request from a Node.js service, with the error envelope handled:

```
const response = await fetch(  
  "https://api.northlight.example/v1/weather/current?station=OSL-01",  
  { headers: { Authorization: `Bearer ${process.env.NORTHLIGHT_KEY}` } },  
)  
const body = await response.json()  
if (!response.ok) throw new Error(`${body.error.code}: ${body.error.message}`)  
console.log(`${body.station}: ${body.temperature} °${body.units.temperature}`)
```

And from Python:

```
import os  
import requests  
  
response = requests.get(  
  "https://api.northlight.example/v1/weather/current",  
  params={"station": "OSL-01"},  
  headers={"Authorization": f"Bearer {os.environ['NORTHLIGHT_KEY']}"},  
  timeout=5,  
)  
body = response.json()  
if not response.ok:  
    raise RuntimeError(f"{body['error']['code']}: {body['error']['message']}")  
print(body["station"], body["temperature"], body["units"]["temperature"])
```

Read the response

Every response carries `observedAt`, the instant the station reported¹, and a `units` object naming the unit of every numeric field. The fields you will use most:

Field	Type	Description
<code>observedAt</code>	string	When the station made the reading, ISO 8601 in UTC.
<code>temperature</code>	number	Air temperature two metres above ground.
<code>windSpeed</code>	number	Ten-minute mean wind speed; gusts are in <code>windGust</code> when the station has it.

Field	Type	Description
<code>windDirection</code>	integer	Degrees clockwise from north, the direction the wind comes from.
<code>conditions</code>	string	One of <code>clear</code> , <code>partly-cloudy</code> , <code>overcast</code> , <code>fog</code> , <code>rain</code> , <code>snow</code> , <code>sleet</code> .
<code>stale</code>	boolean	<code>true</code> when the reading is older than fifteen minutes.

A missing station returns `404` with a machine-readable `code` , listed under [Error codes](#). Branch on the code; the `message` is for humans and may change between releases.

Two response headers matter from the first request on. `X-Limit-Minute` and `X-Limit-Hour` report how many requests remain in the current windows, and `ETag` identifies the reading, so a later request can send `If-None-Match` and receive a `304` that does not count against any window:

```
GET /v1/weather/current?station=OSL-01 HTTP/1.1
Host: api.northlight.example
Authorization: Bearer nlk_live_...
If-None-Match: "obs-OSL-01-1758520800"
```

Stay within the limits

Current conditions are cached at the edge for sixty seconds, so polling faster than that returns the same reading and spends your quota. The windows and the headers that report them are in [Rate limits](#); the short version is below.

Windows

As a new integrator you share one rule: a key has a rolling minute window on `/v1/weather/current` and a rolling hour window across all endpoints, the size of each depends on the plan, and both counters are reported on every response. A request that finds either window empty is answered with `429 rate_limited` and is not counted.

Header	Meaning
<code>X-Limit-Minute</code>	Requests left in the current minute window.
<code>X-Limit-Hour</code>	Requests left in the current hour window.
<code>Retry-After</code>	Seconds to wait, present only on a <code>429</code> .

Windows are rolling, not aligned to the clock: a request made at 12:00:30 leaves the minute window at 12:01:30. A `304` answer to a conditional request counts against neither window, and neither does a `429` or a `5xx`.

A schedule that works for most services:

1. Poll `/v1/weather/current` once a minute per station, offset by a few seconds per station so the requests do not all land at once.
2. Fetch the hourly forecast once an hour, shortly after the model run lands at five past the hour.
3. Send `If-None-Match` on every poll and treat a `304` as "nothing new".
4. On a `429`, wait for `Retry-After`; on a `503`, back off as described in [Backoff](#).

Bursts are counted per key

Two services sharing one key share one window. Give each service its own key so a spike in one cannot starve the other, and so a leaked key can be revoked without stopping everything else.

Go to production

Before switching from the test key to a live one, check that your client:

- reads `units` instead of assuming metric;
- treats `stale: true` as a reason to show the reading's age, not to drop it;
- handles every code in [Error codes](#), including `station_retired`, which arrives one day for any long-running client;
- logs `X-Request-Id` from every response, so support can find the request;
- ignores fields it does not know, since additive changes ship without notice (see [Versioning](#)).

Ready for a live key

A client that passes the list above moves to a live key by changing one environment variable. Nothing else in the request changes between the two kinds of key.

Footnotes

1. Always UTC, always with a `Z` suffix. Convert at the edge of your own system, not in the data layer; the station's IANA time zone is in its [detail record](#). ↩

Endpoint reference

Each second-level section describes one endpoint. The first table of a section is its basic information, which the overview page reads into a summary table; the tables that follow describe parameters and the response. Error responses share one shape and are listed separately under [Error codes](#).

All endpoints are served from `https://api.northlight.example` and answer `GET` only. Requests carry these headers:

Header	Required	Description
<code>Authorization</code>	yes	<code>Bearer</code> followed by a live or test key.
<code>Accept</code>	no	<code>application/json</code> , the only representation; any other value is answered with <code>406</code> .
<code>If-None-Match</code>	no	The <code>ETag</code> of an earlier response; a <code>304</code> means it is still current.
<code>X-Request-Id</code>	no	Your own correlation id, echoed back; one is generated when absent.

List endpoints paginate with a cursor: a response that has more rows carries `page.next`, and passing it as `cursor` returns the next page. Cursors are opaque and expire after ten minutes.

Current conditions

Returns the latest observation from one station.

Method	Path	Auth
GET	<code>/v1/weather/current</code>	Bearer

Parameter	Type	Required	Description
<code>station</code>	string	yes	The station code, such as <code>0SL-01</code> . One station per request.
<code>units</code>	string	no	<code>metric</code> (default) or <code>imperial</code> .

The response is one observation object. Readings older than fifteen minutes carry `stale: true`; a station that has not reported for six hours is answered with `503 upstream_delayed` instead.

Field	Type	Description
<code>station</code>	string	The station code as requested.
<code>observedAt</code>	string	The reading's instant, ISO 8601 in UTC.

Field	Type	Description
<code>temperature</code>	number	Air temperature two metres above ground.
<code>windSpeed</code>	number	Ten-minute mean wind speed.
<code>windGust</code>	number	Highest three-second gust in the same ten minutes; absent without a gust sensor.
<code>windDirection</code>	integer	Degrees clockwise from north.
<code>humidity</code>	integer	Relative humidity in percent.
<code>pressure</code>	number	Pressure reduced to sea level.
<code>precipitation</code>	number	Precipitation in the last ten minutes.
<code>conditions</code>	string	A summary word; see Read the response .
<code>stale</code>	boolean	<code>true</code> when <code>observedAt</code> is more than fifteen minutes old.
<code>units</code>	object	The unit of every numeric field.

Responses are cached at the edge for sixty seconds and carry an `ETag`.

Hourly forecast

Returns up to 48 hourly forecast points for one station.

Method	Path	Auth
GET	<code>/v1/forecast/hourly</code>	Bearer

Parameter	Type	Required	Description
<code>station</code>	string	yes	The station code.
<code>hours</code>	integer	no	1 to 48, default 24. Values above 48 are clamped, not rejected.
<code>units</code>	string	no	<code>metric</code> (default) or <code>imperial</code> .
<code>fields</code>	list	no	<code>temperature</code> <code>wind</code> <code>precipitation</code> - Probability and amount <code>cloudCover</code>

`fields` narrows the response to the named groups; without it every group is returned. The response carries the model run that produced the points and the points themselves:

```
{
  "station": "OSL-01",
  "issuedAt": "2026-09-22T05:05:00Z",
  "points": [
    {
      "validAt": "2026-09-22T07:00:00Z",
      "temperature": 12.1,
      "windSpeed": 4.0,
      "windDirection": 250,
      "precipitationProbability": 0.15,
      "precipitationAmount": 0.0,
      "cloudCover": 0.9
    }
  ],
  "units": {
    "temperature": "C",
    "windSpeed": "m/s",
    "precipitationAmount": "mm"
  }
}
```

Beyond 24 hours the Community plan receives no points; the other plans receive up to 48. See [Plans compared](#).

Daily forecast

Returns one forecast summary per calendar day, up to seven days ahead.

Method		Path		Auth
GET		/v1/forecast/daily		Bearer
Parameter	Type	Required	Description	
station	string	yes	The station code.	
days	integer	no	1 to 7, default 5. Days are calendar days in the station's time zone.	
units	string	no	metric (default) or imperial .	

Each day carries `temperatureMin`, `temperatureMax`, `precipitationTotal`, `precipitationProbability`, the dominant `conditions` word, and `sunrise` and `sunset` in UTC. The first day is today in the station's local time, which may already be partly in the past; its minimum and maximum cover the whole day.

Daily summaries are derived

A day's summary is computed from the hourly points of the same model run, so the two forecast endpoints never disagree about the same run. They can differ when a new run has landed between two requests; compare `issuedAt`.

Observation history

Returns past observations from one station in a time range, oldest first.

Method	Path	Auth
GET	<code>/v1/weather/history</code>	Bearer

Parameter	Type	Required	Description
<code>station</code>	string	yes	The station code.
<code>from</code>	string	yes	Start of the range, ISO 8601, inclusive.
<code>to</code>	string	no	End of the range, exclusive; defaults to now.
<code>step</code>	string	no	<code>10m</code> (default), <code>1h</code> or <code>1d</code> . Coarser steps return means, with min and max for temperature.
<code>limit</code>	integer	no	Rows per page, 1 to 1,000, default 500.
<code>cursor</code>	string	no	<code>page.next</code> from the previous page.

How far back `from` may reach depends on the plan's history retention; a range that starts earlier is answered with `400 range_too_wide` rather than being silently clipped. A range that spans more than 31 days at the `10m` step is rejected the same way; use `1h` or `1d` for long ranges.

```
{
  "station": "OSL-01",
  "step": "1h",
  "rows": [
    {
      "observedAt": "2026-09-21T00:00:00Z",
```

```
    "temperature": 9.8,
    "temperatureMin": 9.1,
    "temperatureMax": 10.4,
    "windSpeed": 2.1,
    "precipitation": 0.0
  }
],
"page": { "next": "eyJvIjoimjAyNi0wOS0yMVQwMTowMDowMFoifQ" },
"units": { "temperature": "C", "windSpeed": "m/s", "precipitation": "mm" }
}
```

Station directory

Lists the stations a key may read, with their location and status.

Method	Path	Auth
GET	/v1/stations	Bearer

Parameter	Type	Required	Description
region	string	no	An ISO 3166-2 code to narrow the list, such as NO-03 .
status	string	no	active , maintenance or retired .
limit	integer	no	Rows per page, 1 to 200, default 50.
cursor	string	no	page.next from the previous page.

Each row carries the code, the site name, coordinates, elevation, status and the instant of the last observation. Retired stations are included only when `status=retired` is asked for, so a client that lists stations sees them disappear rather than fail.

```
{
  "stations": [
    {
      "code": "OSL-01",
      "site": "Oslo",
      "latitude": 59.9423,
      "longitude": 10.7203,
      "elevation": 94,
      "status": "active",
      "lastObservedAt": "2026-09-22T06:00:00Z"
    }
  ]
}
```

```
  },
  {
    "code": "BG0-02",
    "site": "Bergen",
    "latitude": 60.3982,
    "longitude": 5.3241,
    "elevation": 320,
    "status": "maintenance",
    "lastObservedAt": "2026-09-19T14:10:00Z"
  }
],
"page": { "next": null }
}
```

Station detail

Returns one station's metadata, including its sensors and time zone.

Method	Path	Auth
GET	/v1/stations/{code}	Bearer

Parameter	Type	Required	Description
code	string	yes	The station code, in the path.

The detail adds to the directory row the station's IANA time zone, its sensor list with the make and the installation date of each, the date it first reported and, for a retired station, the date it stopped. Metadata changes rarely; cache it for a day.

Field	Type	Description
timezone	string	IANA zone name, such as Europe/Oslo .
sensors	list	- temperature - wind - humidity - pressure - precipitation
since	string	The date the station first reported.
retiredAt	string	Present only for a retired station.

Weather alerts Beta

Lists active severe-weather alerts for a region, relayed from the national service.

Method		Path		Auth
GET		/v1/alerts		Bearer
Parameter	Type	Required	Description	
region	string	yes	An ISO 3166-2 code.	
severity	string	no	yellow , orange or red ; lower levels are excluded.	

Each alert carries event (wind , rain , snow , ice or flood), severity , onset , expires and a headline in the region's language. The endpoint is in beta: its shape may change with thirty days' notice rather than ninety, and it is available on the Business and Enterprise plans.

Not a warning service

Alerts are relayed with a delay of up to five minutes and may be withdrawn. Use them to annotate a forecast, not to trigger safety-critical actions.

Rate limits

Limits protect the stations, which report on fixed schedules, from being asked the same question thousands of times a minute. They are generous for any client that respects the cache, and every response says where you stand, so a well-behaved client never has to guess.

Windows

Every plan shares one rule: a key has a rolling minute window on `/v1/weather/current` and a rolling hour window across all endpoints, the size of each depends on the plan, and both counters are reported on every response. A request that finds either window empty is answered with `429 rate_limited` and is not counted.

Header	Meaning
<code>X-Limit-Minute</code>	Requests left in the current minute window.
<code>X-Limit-Hour</code>	Requests left in the current hour window.
<code>Retry-After</code>	Seconds to wait, present only on a <code>429</code> .

Windows are rolling, not aligned to the clock: a request made at 12:00:30 leaves the minute window at 12:01:30. A `304` answer to a conditional request counts against neither window, and neither does a `429` or a `5xx`.

Backoff

On a `429` or a `503`, wait for `Retry-After` when present, otherwise double the delay from one second up to thirty, and stop after five attempts. Add a random fraction of a second to every delay so that several clients that failed together do not retry together.

```
async function withBackoff(request, attempts = 5) {
  let delay = 1000
  for (let attempt = 1; ; attempt++) {
    const response = await request()
    if (response.status !== 429 && response.status !== 503) return response
    if (attempt === attempts) return response
    const retryAfter = Number(response.headers.get("retry-after"))
    const wait =
      retryAfter > 0 ? retryAfter * 1000 : delay + Math.random() * 1000
    await new Promise((resolve) => setTimeout(resolve, wait))
    delay = Math.min(delay * 2, 30_000)
  }
}
```

```
}  
}
```

Caching

Current conditions are cached at the edge for sixty seconds, so a poll faster than that returns the same reading and spends the minute window for nothing. Forecasts are cached until the next model run, and history responses for a day. Every cached response carries an `ETag` ; sending it back as `If-None-Match` turns an unchanged answer into a `304` that costs nothing.

Request costs

Not every request weighs the same against the hour window. Heavier queries cost more, so a client that pages through history with small pages pays more than one that asks for large pages.

Endpoint	Cost	Notes
Current conditions	1	Also counted in the minute window.
Hourly forecast	2	Regardless of <code>hours</code> and <code>fields</code> .
Daily forecast	2	
Observation history	5 per page	A page of 1,000 rows costs the same as one of 10.
Station directory	1 per page	
Station detail	1	Cache it for a day.
Weather alerts	1	Business and Enterprise plans.

`X-Limit-Hour` reports the remaining hour budget in cost units, not in requests; `X-Limit-Minute` counts requests, since only one endpoint has a minute window.

Plans compared

The plans differ in window size, forecast horizon, history retention and support. The table has seven columns, so the paginated outputs put it on a landscape page of its own.

Plan	Stations	Minute window	Hour window	Forecast hours	History	Support
Community	3	60	600	24	7 days	Forum
Team	25	120	6,000	48	90 days	Email, 2 days
Business	250	600	60,000	48	1 year	Email, 1 day
Enterprise	Custom	Custom	Custom	48	Custom	Named contact

Stations counts the distinct stations a project's keys may read, not the number of keys; a project may hold up to ten keys. The minute window applies per key, and the hour window applies to the project as a whole, so it is shared by every key the project holds.

Changing plans

An upgrade takes effect within a minute and widens the current windows in place, so a client that is being limited recovers without waiting for the hour to roll over. A downgrade takes effect at the end of the billing period. History older than the new retention stops being served at that point but is kept for thirty days, so a downgrade made by mistake can be reversed without loss.

Change	Takes effect	Windows	History
Upgrade	Within a minute	Widened in place	Older history becomes available at once
Downgrade	End of the billing period	Narrowed at the next hour window	Kept for thirty days, then trimmed
Cancellation	End of the billing period	Live keys stop; test keys keep working	Exportable for thirty days

Appendix: measuring your own usage

Log the two limit headers with every response and chart the minimum per minute. A flat line near zero means a client is polling faster than the sixty-second cache refreshes; slow it down rather than raising the plan. One log line per response is enough:

```
2026-09-22T06:00:01Z GET /v1/weather/current station=OSL-01 status=200 minute=57
hour=5931 request=req_7f3a9c2e
```

A query over a day of such lines shows where the window runs lowest:

```
SELECT date_trunc('minute', at) AS minute, min(minute_left) AS worst
FROM northlight_requests
WHERE at > now() - interval '1 day'
GROUP BY 1
ORDER BY worst
LIMIT 20;
```

Three patterns account for almost every limit incident:

- **Synchronized polling.** Every instance of a service wakes on the same second. Offset instances by a few seconds each.
- **Retry storms.** A `503` from one station is retried at once by every instance. Use the backoff above; the jitter is not optional.
- **Missing conditional requests.** A poll that ignores `ETag` pays for every unchanged reading. Send `If-None-Match`.

Retired keys keep counting

Requests made with a retired key are rejected but still counted against the project for one hour, so rotate keys during a quiet period and switch every consumer before the old key is revoked.

Error codes

Every error response has the same shape: an `error` object with a `code` from the table below, a human-readable `message` that may change between releases, the `requestId` of the failed request and, for validation errors, a `details` object naming the parameter. Branch on the code, never on the message.

```
{
  "error": {
    "code": "parameter_invalid",
    "message": "hours must be an integer between 1 and 48",
    "requestId": "req_7f3a9c2e",
    "details": { "parameter": "hours", "received": "0" }
  }
}
```

Field	Type	Description
<code>code</code>	string	One of the codes below. Stable; treat unknown codes as errors.
<code>message</code>	string	For humans and logs. Not stable.
<code>requestId</code>	string	Quote it to support. Also sent as the <code>X-Request-Id</code> header.
<code>details</code>	object	Present on <code>400</code> errors; names the parameter and the value seen.

All codes

Code	HTTP	Meaning	Retry
<code>parameter_missing</code>	400	A required parameter is absent.	No
<code>parameter_invalid</code>	400	A parameter has the wrong type or is out of range.	No
<code>range_too_wide</code>	400	A history range is too long for its step or starts before the retention.	No
<code>key_invalid</code>	401	The <code>Authorization</code> header is missing or malformed.	No
<code>key_revoked</code>	401	The key was revoked or rotated.	No
<code>key_scope</code>	403	The key's project does not cover the station.	No
<code>plan_required</code>	403	The endpoint or the horizon needs a higher plan.	No
<code>station_unknown</code>	404	No station has that code.	No

Code	HTTP	Meaning	Retry
<code>not_acceptable</code>	406	<code>Accept</code> names a representation other than JSON.	No
<code>station_retired</code>	410	The station stopped reporting permanently.	No
<code>rate_limited</code>	429	The window is exhausted.	After <code>Retry-After</code>
<code>upstream_delayed</code>	503	The station has not reported in time.	With backoff
<code>maintenance</code>	503	The API is being upgraded; announced in the release notes.	After <code>Retry-After</code>

Retrying

The backoff rule for `rate_limited`, `upstream_delayed` and `maintenance` is in [Rate limits](#). No 4xx other than 429 is worth retrying without changing the request.

Request errors

A 400 describes a request that can be corrected before it is sent again. The `details` object names the parameter and repeats the value the server saw, so nothing has to be parsed out of the message. Validate ranges before sending: `hours` is 1 to 48, `days` 1 to 7, `limit` 1 to 1,000 for history and 1 to 200 for the directory, and `step` is one of `10m`, `1h` and `1d`.

`range_too_wide` is the one 400 that depends on the plan. It is returned when `from` is earlier than the plan's history retention allows, and when a range at the `10m` step spans more than 31 days. The `details` object carries `earliest`, the first instant the plan can serve, so a client can clamp its range and retry once.

Authorization errors

A 401 means the key itself was not accepted. `key_invalid` is almost always a formatting problem: a missing `Bearer` prefix, a key pasted with a trailing newline from a secrets file, or a test key sent to a live station. `key_revoked` follows a rotation or a revocation and is permanent for that key; deploy the new key everywhere before revoking the old one.

A 403 means the key is fine but the request is not. `key_scope` says the station is not in the key's project; add it in the console rather than creating a second project. `plan_required` says the endpoint or the horizon is not in the plan, and its `details.plan` names the lowest plan that allows it.

Availability errors

A `503` is temporary by definition. `upstream_delayed` means the station has been silent for more than six hours; the station directory shows the same in `lastObservedAt`, and the station usually returns within a day. `maintenance` is rare and announced in the [release notes](#) beforehand, and it always carries `Retry-After`.

A client that sorts codes into three groups covers every case:

```
switch (body.error?.code) {
  case "rate_limited":
  case "upstream_delayed":
  case "maintenance":
    return retryLater(response)
  case "station_retired":
    return removeStation(station)
  case "key_revoked":
    return alertOperators("Northlight key revoked")
  default:
    throw new Error(`${body.error.code}: ${body.error.message}`)
}
```

Exercise every branch

The synthetic station `TEST-01` accepts a `simulate` parameter naming any code in the table, so every branch of your error handling can be tested without waiting for the real thing.

Release notes

Changes to the API, newest first. Dates are the day a change reached every region. Additive changes are listed for information; changes that alter an existing shape are announced under [Upcoming changes](#) at least ninety days before they take effect.

Upcoming changes

Change	Announced	Effective	Affects
The <code>windMs</code> and <code>temperatureC</code> aliases leave current conditions; read <code>windSpeed</code> , <code>temperature</code> and <code>units</code> .	2026-09-15	2026-12-15	Current conditions, history
Weather alerts leave beta; <code>headline</code> becomes an object with one entry per language.	2026-09-15	2026-10-15	Weather alerts
Cursor expiry shortens from one hour to ten minutes.	2026-07-02	2026-10-01	Station directory, history

2026-09-15 Latest

- `fields` on the hourly forecast accepts `precipitation`, returning probability and expected amount, and `cloudCover`. → [Hourly forecast](#)
- `X-Limit-Hour` is now sent on every response, not only when the minute window is low.
- The hour window is charged in cost units; existing budgets were converted so that no client pays more than before. → [Request costs](#)
- The weather alerts endpoint opens in beta for the Business and Enterprise plans. → [Weather alerts](#)
- The synthetic station `TEST-01` accepts `simulate` to reproduce any error code. → [Error codes](#)

2026-07-02

- Readings older than fifteen minutes carry `stale: true` instead of being withheld.
- The `station_retired` error replaces a generic `404` for stations that stopped reporting permanently.
- Observation history accepts `step=1d`, returning daily means with the minimum and maximum temperature.
- Station detail returns `timezone`, `sensors` and `since`.

2026-05-20

- Daily forecast endpoint added, up to seven days ahead. → [Daily forecast](#)
- Every cached response carries an `ETag` , and a `304` answer counts against no window. → [Caching](#)
- The station directory paginates with `limit` up to 200 rows and a `page.next` cursor.

2026-04-10

- First public release of `/v1` : current conditions, hourly forecast, observation history and the station directory.
- Community, Team and Business plans; Enterprise by arrangement.
- Test keys and the synthetic station `TEST-01` .

Deprecation policy

A field or endpoint is announced here at least ninety days before it changes shape, and the previous behaviour stays available for that period. Beta endpoints carry a thirty-day notice instead and say so on their reference page.